

# Music Informatics Final Project: Song Transitions

1. **Feature extraction (Matej)** tempo, downbeats, key, energy, spectral flux, segmentation
2. **Tempo match & alignment (Bailin)**: time-stretch song B to A's BPM, align downbeats, find overlap windows
3. **Transition scoring (Arvid)**: score overlaps, pick best, apply crossfade, output the mix
4. **Evaluation & integration (Matei)**: metrics, beat error, key distance, energy smoothness, listening, logging, report

1. **Feature extraction (Matej)** tempo, downbeats, key, energy, spectral flux, segmentation

```
In [ ]: import matplotlib.pyplot as plt
import numpy as np
from pydub import AudioSegment
import librosa, librosa.display
import scipy
import pickle
```

```
In [ ]: # Plot parameters
params = {'legend.fontsize': 'x-large',
          'figure.figsize': (12, 4),
          'axes.labelsize': 'x-large',
          'axes.titlesize': 'x-large',
          'xtick.labelsize': 'x-large',
          'ytick.labelsize': 'x-large'}
plt.rcParams.update(params)
```

```
In [ ]: def load_audio(path, sr=44100):
        """Load audio and return waveform + sample rate."""
        y, sr = librosa.load(path, sr=sr, mono=True)
        return y, sr
```

```
In [ ]: # Load the songs
path1 = 'under_pressure.mp3'
path2 = 'billie_jean.mp3'
path3 = 'get_lucky.mp3'
y_under_pressure, sr = load_audio(path1)
y_billie_jean, sr = load_audio(path2)
y_get_lucky, sr = load_audio(path3)
```

```
In [ ]: def get_tempo_baseline(y, sr):
        """
```

```
Compute the tempo and beat locations. Returns the global tempo over the
song and the location of beats.
```

```
"""
```

```
tempo, beat_frames = librosa.beat.beat_track(y=y, sr=sr)
beat_times = librosa.frames_to_time(beat_frames, sr=sr)
return tempo[0], beat_times
```

```
In [ ]: tempo_lib_1, beats_lib_1 = get_tempo_baseline(y_under_pressure, sr)
tempo_lib_2, beats_lib_2 = get_tempo_baseline(y_billie_jean, sr)
tempo_lib_3, beats_lib_3 = get_tempo_baseline(y_get_lucky, sr)

print(f"Tempo baseline: {tempo_lib_1:.1f} BPM, according to google should be ~112")
print(f"Tempo baseline: {tempo_lib_2:.1f} BPM, according to google should be ~117")
print(f"Tempo baseline: {tempo_lib_3:.1f} BPM, according to google should be ~116")
```

Tempo baseline: 114.8 BPM, according to google should be ~112

Tempo baseline: 117.5 BPM, according to google should be ~117

Tempo baseline: 114.8 BPM, according to google should be ~116


```
In [ ]: def plot_beats_zoom(y, sr, beat_times, tempo, start_time=0.0, duration=5.0,
    """
    Plot song waveform and beat locations for a specified time window.
    """
    # Get the specified segment
    start_sample = int(start_time * sr)
    end_sample = int((start_time + duration) * sr)
    y_segment = y[start_sample:end_sample]

    # Only keep beats within this segment
    mask = (beat_times >= start_time) & (beat_times <= start_time + duration)
    beats_in_window = beat_times[mask]


    # Compute amplitude limits (for nicer plot so it dynamically scales)
    y_max = np.max(np.abs(y_segment))
    y_margin = y_max * margin_ratio
    ymin, ymax = -y_max - y_margin, y_max + y_margin

    # Plot
    plt.figure(figsize=(12, 3))
    librosa.display.waveshow(y_segment, sr=sr, alpha=0.7)
    plt.vlines(beats_in_window - start_time, ymin=ymin, ymax=ymax,
        color='r', linestyle='--', alpha=0.8, label='Beats')
    plt.ylim(ymin, ymax)
    plt.title(f"Beat Tracking (Zoomed) - Tempo ≈ {tempo:.1f} BPM\n"
        f"Window: {start_time:.1f} - {start_time+duration:.1f} s")
    plt.xlabel("Time (s)")
    plt.legend()
    plt.tight_layout()
    plt.show()
```


```
In [ ]: plot_beats_zoom(y_under_pressure, sr, beats_lib_1, tempo_lib_1, start_time=0.0, duration=5.0)
```



```
In [ ]: plot_beats_zoom(y_billie_jean, sr, beats_lib_2, tempo_lib_2, start_time=10,
```



```
In [ ]: plot_beats_zoom(y_get_lucky, sr, beats_lib_3, tempo_lib_3, start_time=5, dur
```

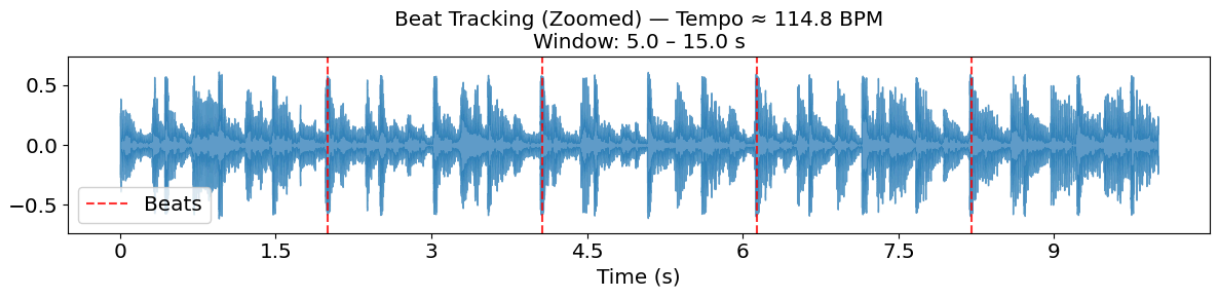


```
In [ ]: def get_downbeats_baseline(beat_times):
    """
    Get downbeats by assuming 4/4 time (same as paper) and the first beat is
    """
    # Naive assumption: first beat is a downbeat. The paper uses a logistic
    # to determine whether the first downbeat is on the first, second, third
    downbeat_times = beat_times[::4]
    return downbeat_times
```

```
In [ ]: downbeats_lib_1 = get_downbeats_baseline(beats_lib_1)
downbeats_lib_2 = get_downbeats_baseline(beats_lib_2)
downbeats_lib_3 = get_downbeats_baseline(beats_lib_3)

print(f"Under Pressure: Found {len(beats_lib_1)} beats and {len(downbeats_lib_1)} downbeats.")
plot_beats_zoom(y_get_lucky, sr, downbeats_lib_3, tempo_lib_3, start_time=5,
```

Under Pressure: Found 457 beats and 115 downbeats.



```
In [ ]: def get_key(y, sr):
        """
        Find the key of the song using chromagram
        """
        # Compute the chromagram
        chroma = librosa.feature.chroma_cqt(y=y, sr=sr)
        chroma_mean = np.mean(chroma, axis=1)
        # Major/minor templates
        maj = np.array([1,0,0,0,1,0,0,1,0,0,0,0])
        min_ = np.array([1,0,0,1,0,0,0,1,0,0,1,0])
        labels = ['C', 'C#', 'D', 'D#', 'E', 'F', 'F#', 'G', 'G#', 'A', 'A#', 'B']
        scores = []
        # Compute a correlation score for each key
        for i in range(12):
            scores.append(np.correlate(np.roll(maj, i), chroma_mean))
            scores.append(np.correlate(np.roll(min_, i), chroma_mean))
        best = np.argmax(scores)
        key_name = labels[best//2] + ('m' if best%2 else '')
        return key_name, float(np.max(scores))
```

```
In [ ]: # Key
key_pap, conf_pap = get_key(y_under_pressure, sr)
print(f"Key baseline Under Pressure: {key_pap} ({conf_pap:.2f}), according to google should be D major

key_pap, conf_pap = get_key(y_billie_jean, sr)
print(f"Key baseline Billie Jean: {key_pap} ({conf_pap:.2f}), according to google should be F sharp minor

key_pap, conf_pap = get_key(y_get_lucky, sr)
print(f"Key baseline Get Lucky: {key_pap} ({conf_pap:.2f}), according to google should be F sharp minor
```

Key baseline Under Pressure: Em (2.20), according to google should be D major

Key baseline Billie Jean: F#m (2.15), according to google should be F sharp minor

Key baseline Get Lucky: Bm (2.58), according to google should be F sharp minor

```
In [ ]: def get_energy_baseline(y, sr, hop_length=512):
        """
        Compute the RMS energy, which represents the loudness.
        """
        rms = librosa.feature.rms(y=y, hop_length=hop_length)[0]
        times = librosa.times_like(rms, sr=sr)
        return times, rms
```

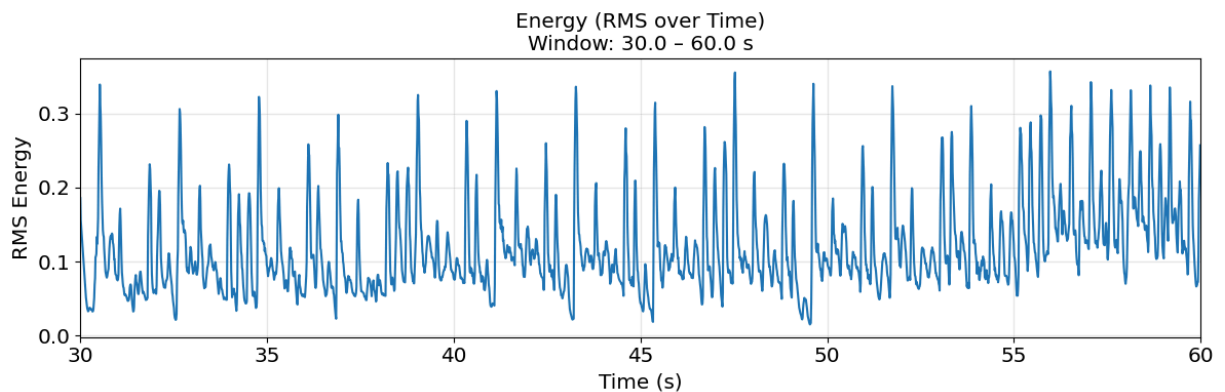
```
In [ ]: def plot_energy(times, energy, title="Energy (RMS over Time)", start_time=0.0,
    """
    Plot energy (RMS) for a specified time window.
    """
    # Calculate end time
    end_time = start_time + duration

    # Get data only from the specified segment
    mask = (times >= start_time) & (times <= end_time)
    times_segment = times[mask]
    energy_segment = energy[mask]

    # Plot
    plt.figure()
    plt.plot(times_segment, energy_segment)
    plt.title(f"{title}\nWindow: {start_time:.1f} - {end_time:.1f} s")
    plt.xlabel("Time (s)")
    plt.ylabel("RMS Energy")
    plt.xlim(start_time, end_time)
    plt.grid(alpha=0.3)
    plt.tight_layout()
    plt.show()
```

```
In [ ]: t_e, rms = get_energy_baseline(y_under_pressure, sr)
plot_energy(t_e, rms, start_time=30)

print(len(t_e))
```



21423

```
In [ ]: def get_spectral_flux_baseline(y, sr):
    """
    Compute the spectral flux using a melspectrogram.
    """
    # Use Melflux ODF
    onset_env = librosa.onset.onset_strength(y=y, sr=sr, feature=librosa.feature.melspectrogram)
    times = librosa.times_like(onset_env, sr=sr)
    return onset_env, times
```

```
In [ ]: def plot_flux(times, flux, title="Spectral Flux", start_time=0.0, duration=30.0,
    """
    Plot spectral flux for a specified time window.
    """
```

```

# Calculate end time
end_time = start_time + duration

# Get data only from the specified segment
mask = (times >= start_time) & (times <= end_time)
times_segment = times[mask]
flux_segment = flux[mask]


# Plot
plt.figure(figsize=(12, 4))
plt.plot(times_segment, flux_segment)
plt.title(f"{title}\nWindow: {start_time:.1f} - {end_time:.1f} s")
plt.xlabel("Time (s)")
plt.ylabel("Flux ( $\Delta$ Spectrum)")
plt.xlim(start_time, end_time)
plt.grid(alpha=0.3)
plt.tight_layout()
plt.show()

```

```

In [ ]: flux_1, flux_times_1 = get_spectral_flux_baseline(y_under_pressure, sr)
plot_flux(flux_times_1, flux_1, start_time=30, duration=10)

```



```

In [ ]: def get_segments_baseline(y, sr, hop_length=512):
    """
    Compute structural novelty curve and segment boundaries (baseline version)
    """
    # Step 1: Compute chroma
    chroma = librosa.feature.chroma_cqt(y=y, sr=sr, hop_length=hop_length)

    # Step 2: Build recurrence matrix (self-similarity)
    R = librosa.segment.recurrence_matrix(chroma, mode='affinity', sym=True)

    # Step 3: Gaussian blur for smoothing
    Rf = scipy.ndimage.gaussian_filter(R, sigma=1.0)

    # Step 4: Build kernel
    L = 16
    kernel = np.outer(np.concatenate([np.ones(L), -np.ones(L)]),
                      np.concatenate([-np.ones(L), np.ones(L)]))

    # Step 5: Compute novelty by 2D convolution along the diagonal
    novelty = np.zeros(Rf.shape[0])
    for i in range(L, Rf.shape[0]-L):

```

```

        novelty[i] = np.sum(Rf[i-L:i+L, i-L:i+L] * kernel)
    novelty = np.maximum(0, novelty) # half-wave rectify

    # Step 6: Peak picking for boundaries
    peaks = librosa.util.peak_pick(novelty, pre_max=8, post_max=8,
                                   pre_avg=8, post_avg=8, delta=0.1, wait=10)
    boundaries = librosa.frames_to_time(peaks, sr=sr, hop_length=hop_length)

    return boundaries, novelty

```

```

In [ ]: def plot_segments(novelty, boundaries, sr, hop_length=512, title="Structural Novelty")
        """
        Plot structural novelty curve and segment boundaries.
        """
        times = librosa.frames_to_time(np.arange(len(novelty)), sr=sr, hop_length=hop_length)

        plt.figure()

        plt.plot(times, novelty, color='teal', alpha=0.7, label='Novelty curve')
        plt.vlines(boundaries, 0, np.max(novelty), color='r', linestyle='--', alpha=0.7)

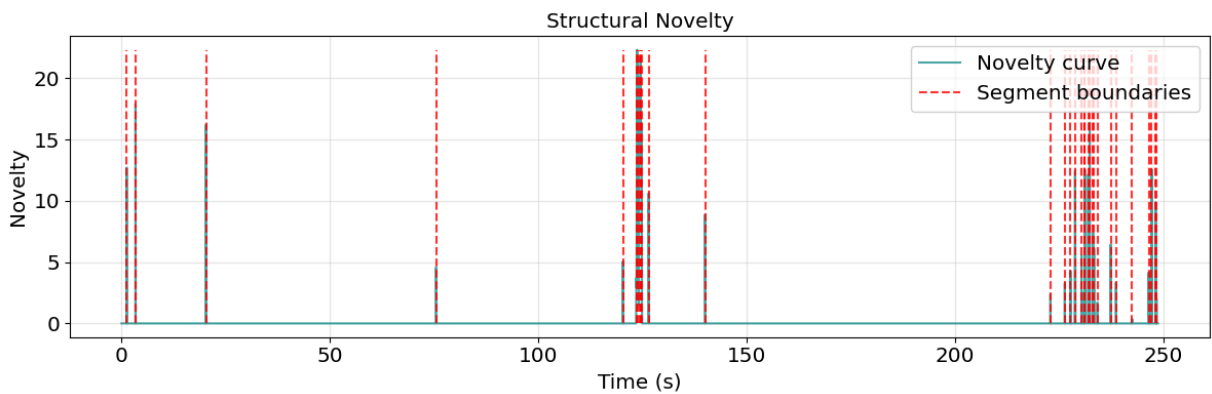
        plt.title(title)
        plt.xlabel("Time (s)")
        plt.ylabel("Novelty")
        plt.legend()
        plt.grid(alpha=0.3)
        plt.tight_layout()
        plt.show()

```

```

In [ ]: boundaries, novelty = get_segments_baseline(y_under_pressure, sr)
        plot_segments(novelty, boundaries, sr)

```

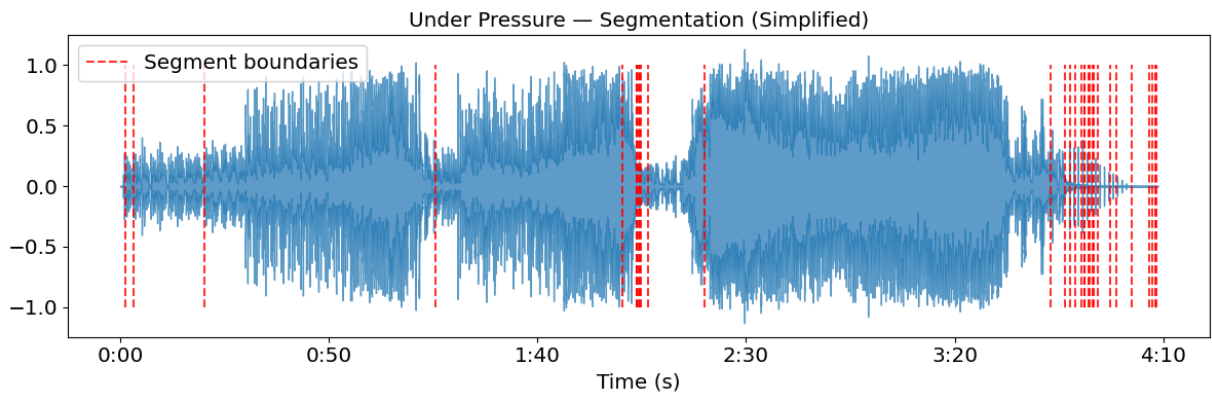


```

In [ ]: def plot_segments_over_waveform(y, sr, boundary_times, title="Structural Segments Over Waveform")
        plt.figure(figsize=(12, 4))
        librosa.display.waveshow(y, sr=sr, alpha=0.7)
        plt.vlines(boundary_times, ymin=-1, ymax=1,
                    color='r', linestyle='--', alpha=0.8, label='Segment boundaries')
        plt.title(title)
        plt.xlabel("Time (s)")
        plt.legend()
        plt.tight_layout()
        plt.show()

```

```
In [ ]: plot_segments_over_waveform(y_under_pressure, sr, boundaries, title="Under Pressure — Segmentation (Simplified)")
```



```
In [ ]: """
Run all the code and all three songs (extend to more if we add them), store
"""
print("Starting feature extraction for all songs...")

# All songs
audio_files = [
    ('Under Pressure', y_under_pressure),
    ('Billie Jean', y_billie_jean),
    ('Get Lucky', y_get_lucky)
]

song_library = []

# Loop through each song and extract all features
for name, y in audio_files:
    print(f"\nProcessing: {name}...")

    # 1. Tempo and Beats
    tempo, beat_times = get_tempo_baseline(y, sr)

    # 2. Downbeats
    downbeat_times = get_downbeats_baseline(beat_times)

    # 3. Key
    key, _ = get_key(y, sr)

    # 4. Energy (RMS)
    energy_values, energy_times = get_energy_baseline(y, sr)

    # 5. Spectral Flux
    flux_times, flux_values = get_spectral_flux_baseline(y, sr)

    # 6. Segmentation
    segment_boundaries, novelty_curve = get_segments_baseline(y, sr)

    # Create the dictionary for this song
    song_data = {
        'name': name,
        'y': y,
        'tempo': tempo,
```



```

        'beat_times': beat_times,
        'downbeat_times': downbeat_times,
        'key': key,
        'energy_values': energy_values,
        'energy_times': energy_times,
        'flux_values': flux_values,
        'flux_times': flux_times,
        'segment_boundaries': segment_boundaries,
        'novelty_curve': novelty_curve
    }

    song_library.append(song_data)

print("\nFeature extraction complete!")

# Save
filename = 'song_library.pkl'
with open(filename, 'wb') as f:
    pickle.dump(song_library, f)

```

Starting feature extraction for all songs...

Processing: Under Pressure...

Processing: Billie Jean...

Processing: Get Lucky...

Feature extraction complete!

```

In [ ]: """
If you have the song_library file, you can use this to load all the data so
Could be useful as the segmentation takes about a minute or two (which is no
"""

filename = 'song_library.pkl'

# Open the file in 'read-binary' (rb) mode and load the data
with open(filename, 'rb') as f:
    loaded_song_library = pickle.load(f)

print(f"Successfully loaded {len(loaded_song_library)} songs from {filename}")

# E.g. use like this
song_a = loaded_song_library[0]
print(f"\nLoaded song name: {song_a['name']}")
print(f"Loaded song tempo: {song_a['tempo']}")


plot_segments_over_waveform(song_a['y'], sr, song_a['segment_boundaries'])

```

Successfully loaded 3 songs from song\_library.pkl

Loaded song name: Under Pressure

Loaded song tempo: 114.84375



2. **Tempo match & alignment (Bailin)**: time-stretch song B to A's BPM, align downbeats, find overlap window

```
In [ ]: import numpy as np
import librosa

"""

def local_tempo(y, sr, center_s, window_s=12.0):

    a = max(0, int((center_s - window_s/2) * sr))
    b = min(len(y), int((center_s + window_s/2) * sr))
    y_win = y[a:b]
    onset = librosa.onset.onset_strength(y=y_win, sr=sr)
    if len(onset) < 4:
        t, _ = librosa.beat.beat_track(y=y_win, sr=sr)
        return float(t)
    tg = librosa.feature.tempogram(onset_envelope=onset, sr=sr)
    tempi = librosa.tempo_frequencies(tg.shape[0], sr=sr)
    curve = tempi[np.argmax(tg, axis=0)]
    return float(np.median(curve))

"""

def local_tempo(y, sr, center_s, window_s=12.0):
    """Median of tempogram peak tempi in a window around center_s"""
    a = max(0, int((center_s - window_s/2) * sr))
    b = min(len(y), int((center_s + window_s/2) * sr))
    if b <= a:
        return float('nan')

    y_win = y[a:b]

    onset = librosa.onset.onset_strength(y=y_win, sr=sr)
    if onset is None or len(onset) < 4:
        t, _ = librosa.beat.beat_track(y=y_win, sr=sr)
        return float(t) if np.isfinite(t) and t > 0 else float('nan')

    tg = librosa.feature.tempogram(onset_envelope=onset, sr=sr)
    if tg.size == 0:
        return float('nan')

    tempi = librosa.tempo_frequencies(tg.shape[0], sr=sr)
```

```

curve = tempi[np.argmax(tg, axis=0)]
val = float(np.median(curve))
return val if np.isfinite(val) and val > 0 else float('nan')

def match_tempo_B_to_A_local(yA, yB, sr, tA_center_s, tB_center_s=0.0, clamp
#Stretch B so its local tempo near tB_center_s matches A near tA_center_s
tempoA = local_tempo(yA, sr, center_s=tA_center_s)
tempoB = local_tempo(yB, sr, center_s=tB_center_s)
raw = tempoA / max(tempoB, 1e-6)
# optional comfort clamp to reduce artifacts
if clamp:
    raw = float(np.clip(raw, clamp[0], clamp[1]))

if tempoB <= 0:
    print(f"***Warning: Calculated tempoB is {tempoB:.2f}. Defaulting to 1.0")
    raw = 1.0
else:
    raw = float(tempoA) / tempoB
    if not np.isfinite(raw) or raw <= 0:
        print(f"***Warning: Invalid stretch rate {raw:.2f}. Defaulting to 1.0")
        raw = 1.0

return librosa.effects.time_stretch(yB, rate=raw), raw, tempoA, tempoB

```

```

In [ ]: def start_time_A_20th_last_beat(beatsA, k=20):
        """
        beatsA: 1D np.array of beat times (seconds) for Track A
        k:      how many beats before the end (default 20)
        """
        if beatsA is None or len(beatsA) == 0:
            raise ValueError("No beats found for Track A.")
        if len(beatsA) <= k:
            return float(beatsA[0])
        return float(beatsA[-k])

def nearest_downbeat_at_or_before(t, downbeats, tol=0.08):
    """
    Return nearest downbeat <= t within 'tol' seconds; else None.
    """
    if downbeats is None or len(downbeats) == 0:
        return None
    idx = np.searchsorted(downbeats, t, side="right") - 1
    if idx < 0:
        return None
    cand = float(downbeats[idx])
    return cand if (t - cand) <= tol else None

def start_time_A_20th_last_with_downbeat(beatsA, downbeatsA=None, k=20, snap
    """
    Pick the k-th last beat as the transition start and, if a downbeat is
    at/just before that moment (within snap_tol), snap to that downbeat.
    """
    tA = start_time_A_20th_last_beat(beatsA, k=k)
    db = nearest_downbeat_at_or_before(tA, downbeatsA, tol=snap_tol)
    return float(db) if db is not None else float(tA)

```

```

def split_track_A_for_transition(yA, sr, start_time_A, overlap_beats, beatsA)
    """
    Split Track A around the chosen transition start.

    Returns:
        A_pre      : samples before the transition start
        A_overlap   : samples during the overlap window (length = overlap_beats)
        A_after     : samples after the overlap window
        end_time_A  : end time (seconds) of the overlap window on A
    """
    # find beat index at/after start_time_A
    i0 = np.searchsorted(beatsA, start_time_A)
    # end index = i0 + overlap_beats (clamped)
    i1 = min(i0 + overlap_beats, len(beatsA) - 1)

    s0 = int(round(start_time_A * sr))
    s1 = int(round(float(beatsA[i1]) * sr))
    s0 = max(0, min(s0, len(yA)))
    s1 = max(s0, min(s1, len(yA)))

    A_pre      = yA[:s0]
    A_overlap   = yA[s0:s1]
    A_after     = yA[s1:]
    return A_pre, A_overlap, A_after, float(beatsA[i1])

```

```

In [ ]: # ----- ENERGY MATCH + EQUAL-POWER CROSSFADE -----

def _rms(x: np.ndarray) -> float:
    return float(np.sqrt(np.mean(np.square(x)) + 1e-12))

def loudness_match_gain(A_overlap: np.ndarray, B_overlap: np.ndarray, sr: int)
    """
    Compute a gain to scale B so its loudness near the seam matches A.
    Uses the last 'ref' window of A and the first 'ref' window of B.
    """
    if len(A_overlap) == 0 or len(B_overlap) == 0:
        return 1.0
    ref_samples = int(min(max_ref_seconds, 0.5 * len(A_overlap) / sr) * sr)
    ref_samples = max(ref_samples, int(0.25 * sr)) # ensure >= ~250 ms

    A_ref = A_overlap[-ref_samples:] if len(A_overlap) >= ref_samples else A
    B_ref = B_overlap[:ref_samples] if len(B_overlap) >= ref_samples else B

    rms_A = _rms(A_ref)
    rms_B = _rms(B_ref)
    return (rms_A / max(rms_B, 1e-9))

def equal_power_crossfade(A_overlap: np.ndarray, B_overlap: np.ndarray) -> np.ndarray
    """
    Apply an equal-power crossfade (sin/cos law) over the overlap.
    Inputs must be the same length; returns the mixed overlap.
    """
    n = min(len(A_overlap), len(B_overlap))
    if n == 0:

```

```

        return np.zeros(0, dtype=np.float32)
    # trim/pad to same length
    a = A_overlap[:n]
    b = B_overlap[:n]
    t = np.linspace(0.0, 1.0, n, dtype=np.float32)
    gA = np.cos(0.5 * np.pi * t)  # 1 -> 0
    gB = np.sin(0.5 * np.pi * t)  # 0 -> 1
    return gA * a + gB * b

def energy_match_and_crossfade(A_overlap: np.ndarray, B_overlap: np.ndarray,
    """
    Scales B to match A's loudness near the seam, then equal-power crossfade
    """
    # 1) loudness match B to A (near the join)
    gain_B = loudness_match_gain(A_overlap, B_overlap, sr)
    B_matched = B_overlap * gain_B
    # 2) equal-power crossfade
    return equal_power_crossfade(A_overlap, B_matched)

```

```

In [ ]: # ----- KEY / PITCH HELPERS -----

_KEY_NAMES = ['C', 'C#', 'D', 'D#', 'E', 'F', 'F#', 'G', 'G#', 'A', 'A#', 'B']

_KEY_NAMES_MAP = {name: i for i, name in enumerate(_KEY_NAMES)}

def format_key(tonic: int, mode: str) -> str:
    return f" {_KEY_NAMES[int(tonic) % 12]} {mode}"

def semitone_diff_to_match(
    tonic_src: int, mode_src: str,
    tonic_tgt: int, mode_tgt: str,
    prefer_same_mode: bool = True
) -> int:
    """
    Smallest signed semitone shift to move src key center to tgt key center.
    If prefer_same_mode and modes match, just align tonics directly.
    If modes differ, still align pitch class (keeps harmonic proximity simple)
    Returned range is [-6, +6].
    """
    diff = (int(tonic_tgt) - int(tonic_src)) % 12
    if diff > 6: diff -= 12
    return int(diff)

def maybe_pitch_shift_B_to_A(
    yB: np.ndarray, sr: int,
    keyA_tonic: int, keyA_mode: str,
    keyB_tonic: int, keyB_mode: str,
    max_semitones: int = 3,          # clamp to avoid artifacts
    prefer_same_mode: bool = True
):
    """
    Compute nearest semitone shift to align B's key center to A's.
    Shift only if |n_steps| <= max_semitones.
    Returns: (yB_shifted, n_steps_applied, keyA_str, keyB_str)
    """
    keyA_str = format_key(keyA_tonic, keyA_mode)

```

```

keyB_str = format_key(keyB_tonic, keyB_mode)

n_steps = semitone_diff_to_match(keyB_tonic, keyB_mode, keyA_tonic, keyA_mode)
if abs(n_steps) > max_semitones:
    # too large: skip to avoid quality loss
    return yB, 0, keyA_str, keyB_str

if n_steps == 0:
    return yB, 0, keyA_str, keyB_str

yB_shifted = librosa.effects.pitch_shift(yB, sr=sr, n_steps=float(n_steps))
return yB_shifted, int(n_steps), keyA_str, keyB_str

```

```

In [ ]: def loudness_match_gain_A_to_B(A_overlap: np.ndarray, B_overlap: np.ndarray,
    """
    Compute a gain to scale **A** so its loudness near the seam matches **B**
    Uses the last window of A and the first window of B. B is left untouched
    """
    if len(A_overlap) == 0 or len(B_overlap) == 0:
        return 1.0

    ref_samples = int(min(max_ref_seconds, 0.5 * len(A_overlap) / sr) * sr)
    ref_samples = max(ref_samples, int(0.25 * sr))

    A_ref = A_overlap[-ref_samples:] if len(A_overlap) >= ref_samples else A_overlap
    B_ref = B_overlap[:ref_samples] if len(B_overlap) >= ref_samples else B_overlap

    try:
        rms_A = _rms(A_ref)
        rms_B = _rms(B_ref)
    except NameError:
        def _rms(x: np.ndarray) -> float:
            return float(np.sqrt(np.mean(np.square(x)) + 1e-12))
        rms_A = _rms(A_ref)
        rms_B = _rms(B_ref)

    return (rms_B / max(rms_A, 1e-9))

def maybe_pitch_shift_A_to_B(
    yA: np.ndarray, sr: int,
    keyA_tonic: int, keyA_mode: str,
    keyB_tonic: int, keyB_mode: str,
    max_semitones: int = 3,
    prefer_same_mode: bool = True
):
    n_steps = semitone_diff_to_match(keyA_tonic, keyA_mode, keyB_tonic, keyB_mode)
    if abs(n_steps) == 0 or abs(n_steps) > max_semitones:
        return yA, 0
    yA_shifted = librosa.effects.pitch_shift(yA, sr=sr, n_steps=float(n_steps))
    return yA_shifted, int(n_steps)

```

```

In [ ]: # 1. MATCH TEMPO
# ratio = get_tempo_baseline(y_A) / get_tempo_baseline(y_B)
# librosa.effects.time_stretch(y_A, ratio)

```

```

# 2. START NEW SONG AT 20TH LAST BEAT OF FIRST SONG
# 20

# 3. MATCH ENERGY
# crossfade A into B

# 4. MATCH PITCH
# (EITHER LIBROSA)
# https://librosa.org/doc/0.11.0/generated/librosa.effects.pitch\_shift.html

# (OR TORCH)
# https://pypi.org/project/torch-pitch-shift/#
# https://kentonishi.com/torch-pitch-shift/usage/

```

3. **Transition scoring (Arvid):** generate outputs, evaluation.

```

In [ ]: # Compare transition positions between A and B
        # Rate each position by beat offset/key/energy
        # Pick best position, match loudness, apply crossfade, save to WAV/MP3

```

```

In [ ]: def parse_key_string(key_str):
        """
        Parses a key string into tonic (INT) and mode (STR)
        """
        key_str = key_str.strip()
        if key_str.endswith('m'):
            tonic_str = key_str[:-1]
            mode = 'minor'
        else:
            tonic_str = key_str
            mode = 'major'

        tonic_int = _KEY_NAMES_MAP.get(tonic_str, 0)
        return tonic_int, mode

```

```

In [ ]: def candidate_transitions(songA, songB, sr, overlap_beats=8, num_downbeats=5):
        """
        Generate candidate transition positions between A and B.
        Prioritizes structural intros/outros if segmentation is available.
        Falls back to first/last N downbeats if structural segments are empty.

        Returns: List of tuples (tA_start, tB_start) of transition times
        """
        downbeats_A = songA['downbeat_times']
        downbeats_B = songB['downbeat_times']

        candidates_A = []
        candidates_B = []

        if 'segment_boundaries' in songA and len(songA['segment_boundaries']) > 0:
            # The last segment starts at the last boundary time
            outro_start_time = songA['segment_boundaries'][-1]

```

```

        # Get all downbeats that fall within this last segment
        candidates_A = downbeats_A[downbeats_A >= outro_start_time]

    if 'segment_boundaries' in songB and len(songB['segment_boundaries']) > 0:
        # The first segment ends at the first boundary time
        intro_end_time = songB['segment_boundaries'][0]
        # Get all downbeats from the start of the song up to this first boundary time
        candidates_B = downbeats_B[downbeats_B < intro_end_time]

    # If we couldn't find any downbeats in Song A's outro, use the last N downbeats
    if len(candidates_A) == 0:
        print(f"Warning: No structural outro downbeats found for {songA['name']}")
        num_A = min(num_downbeats, len(downbeats_A))
        candidates_A = downbeats_A[-num_A:]

    if len(candidates_B) == 0:
        print(f"Warning: No structural intro downbeats found for {songB['name']}")
        num_B = min(num_downbeats, len(downbeats_B))
        candidates_B = downbeats_B[:num_B]

    # Each pair (tA, tB) is a possible transition point
    candidates = []
    for tA in candidates_A:
        for tB in candidates_B:
            candidates.append((tA, tB))

    return candidates

```

```

In [ ]: def rate_transition(songA, songB, tA, tB, sr, overlap_beats=16, silence_threshold=0.05):
    """
    Rates a transition based on key, energy, and rhythmic similarity (DTW).
    """

    # Get beat times for both songs
    beatsA = songA['beat_times']
    beatsB = songB['beat_times']

    # Find the index of the downbeat (tA and tB)
    tA_idx = np.argmin(abs(beatsA - tA))
    tB_idx = np.argmin(abs(beatsB - tB))

    # Key distance
    if songA['key'] == songB['key']:
        key_diff = 0
    else:
        key_diff = 1

    # Find the end of the overlap window
    tA_end_idx = min(tA_idx + overlap_beats, len(beatsA) - 1)
    tB_end_idx = min(tB_idx + overlap_beats, len(beatsB) - 1)

    # Get the ODF (flux) times
    flux_times_A = songA['flux_times']
    flux_times_B = songB['flux_times']

```



```

# Find the start/end times of the overlap in seconds
tA_start_s = beatsA[tA_idx]
tA_end_s   = beatsA[tA_end_idx]
tB_start_s = beatsB[tB_idx]
tB_end_s   = beatsB[tB_end_idx]

# Find the ODF values within that time window
odf_A_indices = (flux_times_A >= tA_start_s) & (flux_times_A < tA_end_s)
odf_B_indices = (flux_times_B >= tB_start_s) & (flux_times_B < tB_end_s)

odf_A_segment = songA['flux_values'][odf_A_indices]
odf_B_segment = songB['flux_values'][odf_B_indices]

if len(odf_A_segment) == 0 or len(odf_B_segment) == 0:
    return float('inf')

# Normalize the ODF segments
odf_A_norm = (odf_A_segment - np.mean(odf_A_segment)) / np.std(odf_A_seg
odf_B_norm = (odf_B_segment - np.mean(odf_B_segment)) / np.std(odf_B_seg

# Calculate rhythmic distance using DTW
D, wp = librosa.sequence.dtw(X=odf_A_norm, Y=odf_B_norm, metric='euclidean')
rhythmic_diff = D[wp[-1, 0], wp[-1, 1]] / len(wp)

# Final Weighted Score
score = (key_diff * 2) + (rhythmic_diff * 1.0)

return score

```

```

In [ ]: def find_best_transition(songA, songB, sr, num_downbeats=5):
        """
        Evaluate candidate transitions and return the best one.
        """
        candidates = candidate_transitions(songA, songB, sr, num_downbeats=num_c

        scored = []
        for (tA, tB) in candidates:
            score = rate_transition(songA, songB, tA, tB, sr)
            scored.append((score, tA, tB))

        # Sort candidates by score
        scored.sort(key=lambda x: x[0])

        # Choose the best one
        best_score, best_tA, best_tB = scored[0]

        return best_tA, best_tB

```

```

In [ ]: def make_transition(songA, songB, sr, overlap_beats=8):
        print(f"Starting transition A to B: {songA['name']} → {songB['name']}")

        yA = np.asarray(songA['y'])
        yB = np.asarray(songB['y'])

        # Anchor points

```

```

best_tA, best_tB = find_best_transition(songA, songB, sr)
print(f"Found best anchor points: tA={best_tA:.2f}s, tB={best_tB:.2f}s")

# Get beat times
beatsA = songA['beat_times']
beatsB = songB['beat_times']

# A
iA0 = int(np.searchsorted(beatsA, best_tA))
iA1 = min(iA0 + overlap_beats, len(beatsA) - 1)
tA0 = float(best_tA)
tA1 = float(beatsA[iA1])

# B
iB0 = int(np.searchsorted(beatsB, best_tB))
iB1 = min(iB0 + overlap_beats, len(beatsB) - 1)
tB0 = float(best_tB)
tB1 = float(beatsB[iB1])

# Slice audio
sA0, sA1 = int(round(tA0 * sr)), int(round(tA1 * sr))
sB0, sB1 = int(round(tB0 * sr)), int(round(tB1 * sr))
A_pre = yA[:sA0]
A_overlap0 = yA[sA0:sA1]
B_overlap = yB[sB0:sB1]

if len(A_overlap0) == 0 or len(B_overlap) == 0:
    print("Warning: empty overlap window")
    return np.concatenate([A_pre, yB[sB0:]])

# Tempo-match A's overlap to B's overlap
dur_A = len(A_overlap0) / sr
dur_B = len(B_overlap) / sr

rate = (dur_A / max(dur_B, 1e-9))
A_overlap = librosa.effects.time_stretch(A_overlap0, rate=float(rate))

# Trim or pad A_overlap to match B_overlap length
n = len(B_overlap)
if len(A_overlap) > n:
    A_overlap = A_overlap[:n]
elif len(A_overlap) < n:
    A_overlap = np.pad(A_overlap, (0, n - len(A_overlap)))

# Pitch-shift A toward B
keyA_tonic, keyA_mode = parse_key_string(songA['key'])
keyB_tonic, keyB_mode = parse_key_string(songB['key'])
A_overlap, n_steps = maybe_pitch_shift_A_to_B(
    A_overlap, sr, keyA_tonic, keyA_mode, keyB_tonic, keyB_mode
)
if n_steps != 0:
    print(f"Pitch-shifted A by {n_steps:+d} semitones toward B during ov

# Loudness-match A to B
gain_A = loudness_match_gain_A_to_B(A_overlap, B_overlap, sr)
A_overlap_matched = A_overlap * gain_A

```

```

# Equal-power crossfade
mixed_overlap = equal_power_crossfade(A_overlap_matched, B_overlap)

# Final assembly
B_post = yB[sB1:]
full_mix = np.concatenate([A_pre, mixed_overlap, B_post])

# Normalize to avoid clipping
peak = np.max(np.abs(full_mix))
if peak > 0:
    full_mix = 0.98 * (full_mix / peak)

print("Transition complete (A morphed into B; B unchanged outside overlap)")
return full_mix

```

```

In [ ]: from scipy.io.wavfile import write

def save_mix(y, sr, filename="mix.wav"):
    y16 = np.int16(y / np.max(np.abs(y)) * 32767)
    write(filename, sr, y16)
    print(f"Saved mix to {filename}")

song_A = loaded_song_library[0]
song_B = loaded_song_library[1]

mix = make_transition(song_A, song_B, sr, overlap_beats=16)
save_mix(mix, sr, "final_output.wav")

```

Starting transition A to B: Under Pressure → Billie Jean  
Warning: No structural outro downbeats found for Under Pressure. Using last 5 downbeats as fallback.  
Found best anchor points: tA=232.68s, tB=1.24s  
Pitch-shifted A by +2 semitones toward B during overlap.  
Transition complete (A morphed into B; B unchanged outside overlap).  
Saved mix to final\_output.wav

#### 4. Evaluation & integration (Matei): evaluation, logging, report

In [ ]:

Resources: <https://asmp-eurasipjournals.springeropen.com/articles/10.1186/s13636-018-0134-8>